

📄 Quiz name

bigdata final 2022

📅 Date played

Oct 21, 2022 1:13 PM

👤 Hosted by

Dr. Chantana CHANTRAPORNCHAI

Accuracy

84%



Class accuracy

56%



Graded Questions

50



Correct

42



Incorrect

8



Avg. time taken

29.4s



Replay this quiz with your child

Play now

Find similar quizzes

Detailed Report

☰ 50 questions

Question 1 **Correct**

🕒 44 secs

🏆 1 points

Question Type: Multiple-choice

What is correct about RDD ? (Select all that apply)

Papon CHOONHAKLAI's Response

It supports data partitioning and distributed through cluster.

data frame can be converted to RDD with python supported.

CSV and JSON can be accepted.

Correct answer

It supports data partitioning and distributed through cluster.

data frame can be converted to RDD with python supported.

CSV and JSON can be accepted.

Question 2 **Incorrect**

🕒 26 secs

🏆 0 points

Question Type: Fill-in-the-Blank

Two types of RDD operations are,

Papon CHOONHAKLAI's Response

reduce,count

Correct answer

transformation and action

Question 3 Correct

🕒 25 secs 📊 1 points

Question Type: Multiple-choice

What is the RDD operation that gives outputs? (select all that apply)

Papon CHOONHAKLAI's Response

Reduce Count

Correct answer

Reduce Count

Question 4 Incorrect

🕒 39 secs 📊 0 points

Question Type: Multiple-choice

What is true about RDD Persistence? (Select all that apply)

Papon CHOONHAKLAI's Response

It caches results in memory only as it it in-memory processing.

The automatic recomputation is done if RDD partition is lost.

Correct answer

It has cache to speedup future operations.

The automatic recomputation is done if RDD partition is lost.

Question 5 Correct

🕒 43 secs 📊 1 points

Question Type: Multiple-choice

What is true about Spark and Hadoop mapreduce? (Select all that apply)

Papon CHOONHAKLAI's Response

Spark has faster performance due to its in-memory processing feature.

Spark can run on top of HDFS or not.

Correct answer

Spark has faster performance due to its in-memory processing feature.

Spark can run on top of HDFS or not.

Question 6 Correct

🕒 12 secs 📊 1 points

Question Type: Multiple-choice

when using jupyter notebook, the command

ssh -N -L p1:localhost:p2 hadoop@10.3.133.231 what is p1, p2?

Papon CHOONHAKLAI's Response

p1 is local port


p2 is remote machine port


Correct answer

p1 is local port

p2 is remote machine port

Question 7 **Correct**

 39 secs

 1 points

Question Type: Multiple-choice

what cannot we do in one command for graphframe? (Select all that apply)


Papon CHOONHAKLAI's Response


find subgraph find depthfirst search

Correct answer

find subgraph find depthfirst search

Question 8 **Correct**

 37 secs

 1 points

Question Type: Multiple-choice

Kafka broker has functions: (Select all that apply)


Papon CHOONHAKLAI's Response

message exchange center Monitor the incoming message buffer the message

Correct answer

message exchange center Monitor the incoming message buffer the message

Question 9 **Correct**

 21 secs

 1 points

Question Type: Multiple-choice

what is the important element we need when running kafka?


Papon CHOONHAKLAI's Response

zookeeper

Correct answer

zookeeper

Question 10 **Correct**

 7 secs

 1 points

Question Type: Multiple-choice

what is the default port for kafka broker?

Papon CHOONHAKLAI's Response

9092

Correct answer

Question 11 Correct 39 secs  1 points

Question Type: Multiple-choice

What is correct about data storage in HBase? (select all that apply)

Papon CHOONHAKLAI's Response

Table is divided by row key. Rows are split into region.

Correct answer

Table is divided by row key. Rows are split into region.

Question 12 Correct 45 secs  1 points

Question Type: Multiple-choice



What is correct about the design criteria for HBase? (Select all that apply)

Papon CHOONHAKLAI's Response

The redundancy is allowed to increase parallelism. Empty column will not be stored.

Correct answer

The redundancy is allowed to increase parallelism. Empty column will not be stored.

Question 13 Incorrect 14 secs  0 points

Question Type: Multiple-choice

Hive usesfor execution engine.

Papon CHOONHAKLAI's Response

Zookeeper

Correct answer

Ted

Question 14 Correct 33 secs  1 points

Question Type: Multiple-choice

HiveServer2 is (select all that apply)

Papon CHOONHAKLAI's Response

for handing multiuser session. connected through jdbc/odbc

Correct answer

for handing multiuser session. connected through jdbc/odbc

Question 15 Correct 27 secs  1 points

Question Type: Multiple-choice

What is correct about Pig? (Select all that apply)

Papon CHOONHAKLAI's Response

It is a kind of data-flow language.


The script is transformed into several Mapreduce tasks.

Correct answer

It is a kind of data-flow language.

The script is transformed into several Mapreduce tasks.

Question 16 Correct

 18 secs  1 points

Question Type: Multiple-choice

The following command in Pig triggers mapreduce job. (Select all that apply)



Papon CHOONHAKLAI's Response

dump store

Correct answer

dump store

Question 17 Correct

 33 secs  1 points

Question Type: Multiple-choice

What are not important elements in hbase-site.xml?

Papon CHOONHAKLAI's Response

dfs replication

Correct answer

dfs replication

Question 18 Correct

 9 secs  1 points

Question Type: Fill-in-the-Blank

what is the default port for kibana?

Papon CHOONHAKLAI's Response

5601

Correct answer

5601

Question 19 Correct

 45 secs  1 points

Question Type: Multiple-choice

what is correct about elasticsearch? (Select all that apply)

Papon CHOONHAKLAI's Response

it supports the full text Index is equivalent to we query the document based

search.

document

on json.

Correct answer

it supports the full text
search.

Index is equivalent to
document

we query the document based
on json.

Question 20 Correct

🕒 36 secs 🏆 1 points

Question Type: Multiple-choice

what is correct about logstash ? (select all that apply)

Papon CHOONHAKLAI's Response

it can handle
multiple data
source.

we use logstash to filter data before
entering elasticsearch to reduce the
datasize

one filebeat instance
can handle multiple
files.

Correct answer

it can handle
multiple data
source.

we use logstash to filter data before
entering elasticsearch to reduce the
datasize

one filebeat instance
can handle multiple
files.

Question 21 Correct

🕒 35 secs 🏆 1 points

Question Type: Multiple-choice

Logstash can output to more than just elasticsearch.

Papon CHOONHAKLAI's Response

True

Correct answer

True

Question 22 Correct

🕒 35 secs 🏆 1 points

Question Type: Fill-in-the-Blank

what is the keyword for output of logstash to output to screen?

Papon CHOONHAKLAI's Response

rubydebug

Correct answer

stdout

rubydebug

Question 23 Correct

🕒 24 secs 🏆 5 points

Question Type: Multiple-choice

which command can be used to test the existence of Elasticsearch process on terminal?

Papon CHOONHAKLAI's Response

curl

Correct answer

curl

Question 24 Correct

🕒 27 secs 📊 5 points

Question Type: Multiple-choice

if elasticsearch gets killed during the starting up, what is the likely cause? (select all that apply)

Papon CHOONHAKLAI's Response

jvm heap size is too large for available memory.

Correct answer

jvm heap size is too large for available memory.

Question 25 Correct

🕒 17 secs 📊 5 points

Question Type: Multiple-choice

logstash can acceptnumber of input sources.

Papon CHOONHAKLAI's Response

multiple

Correct answer

multiple

Question 26 Correct

🕒 12 secs 📊 5 points

Question Type: Multiple-choice

.....is a feature for logstash used to filter incoming data.

Papon CHOONHAKLAI's Response

grok

Correct answer

grok

Question 27 Correct

🕒 33 secs 📊 5 points

Question Type: Multiple-choice

index is equivalent to in mysql.

Papon CHOONHAKLAI's Response

database

Correct answer

database

Question 28 **Correct**

🕒 30 secs 📊 5 points

Question Type: Multiple-choice

what is advantage of Parquet over pandas? (select all that apply)

Papon CHOONHAKLAI's Response

data compression have embedded schema

Correct answer

data compression have embedded schema

Question 29 **Correct**

🕒 55 secs 📊 5 points

Question Type: Multiple-choice

Apache parquet (select all that apply)

Papon CHOONHAKLAI's Response

is a column data store is stored in a key-value format. can be converted from pandas

Correct answer

is a column data store is stored in a key-value format. can be converted from pandas

Question 30 **Correct**

🕒 7 secs 📊 5 points

Question Type: Multiple-choice

Apache parquet data store is saved in a single file.

Papon CHOONHAKLAI's Response

False

Correct answer

False

Question 31 **Incorrect**

🕒 28 secs 📊 0 points

Question Type: Multiple-choice

What is advantage of dask over pandas? (select all that apply)

Papon CHOONHAKLAI's Response

can run on GPU can run on cluster can read any size data frame.

Correct answer

can read multiple files at the same time. can run on GPU can run on cluster can read any size data frame.

Question 32 **Correct**

🕒 25 secs 📊 5 points

Question Type: Multiple-choice

```
result = df.attribute1.mean()
```

what is the next line in order to show the value of mean of attribute1?



Papon CHOONHAKLAI's Response

```
result.compute()
```

Correct answer

```
result.compute()
```

Question 33 **Correct**

 47 secs  5 points

Question Type: Multiple-choice

when the user browses the visualization and plays interactive, what affects the response time to the user? select all that apply.


Papon CHOONHAKLAI's Response

server CPU speed bus speed between memory and cpu internet bandwidth

Correct answer

server CPU speed bus speed between memory and cpu internet bandwidth

Question 34 **Incorrect**

 32 secs  0 points

Question Type: Multiple-choice

what is scheduled into each SM?


Papon CHOONHAKLAI's Response

thread

Correct answer

thread block

Question 35 **Correct**

 22 secs  5 points

Question Type: Multiple-choice

Increasing more cores always increases performance.

Papon CHOONHAKLAI's Response

False

Correct answer

False

Question 36 **Incorrect**

 16 secs  0 points

Question Type: Multiple-choice

GPU centric data movement tries to (select all that apply)

Papon CHOONHAKLAI's Response

Move data between GPUs among nodes via RDMA

Move data between host and device

Correct answer

Move data between GPUs among nodes via RDMA

Move data between GPUs in a host using NVLink

Question 37 Correct

🕒 51 secs 📊 5 points

Question Type: Multiple-choice

What is pinned memory?

Papon CHOONHAKLAI's Response

Unpageable memory

Correct answer

Unpageable memory

Question 38 Correct

🕒 28 secs 📊 5 points

Question Type: Multiple-choice

NVTabular dataset supports.....(Select all that apply)

Papon CHOONHAKLAI's Response

Dask Pandas DataFrame Parquet file

Correct answer

Dask Pandas DataFrame Parquet file

Question 39 Correct

🕒 3 secs 📊 5 points

Question Type: Multiple-choice

NVTabular can read multiple files at one command.

Papon CHOONHAKLAI's Response

True

Correct answer

True

Question 40 Correct

🕒 36 secs 📊 5 points

Question Type: Multiple-choice

To initialize workflow, parameters needed are(select all that apply)



Papon CHOONHAKLAI's Response

column name categorial columns continuous value columns

Correct answer

column name categorial columns continuous value columns

Question 41 **Correct**

 31 secs  5 points

Question Type: Multiple-choice

What are the types of map operations? select all that apply.

Papon CHOONHAKLAI's Response

x=x+1 read (x)

Correct answer

x=x+1 read (x)

Question 42 **Incorrect**

 51 secs  0 points

Question Type: Multiple-choice

What are the ops in NVTabular for feature engineering and preprocessing ? select all that apply.

Papon CHOONHAKLAI's Response

normalize zerofill

Correct answer

normalize log zerofill

Question 43 **Correct**

 50 secs  5 points

Question Type: Multiple-choice

In which case, parquet is faster than csv format. Select all that apply.

Papon CHOONHAKLAI's Response

select one column

Correct answer

select one column

Question 44 **Correct**

 23 secs  5 points

Question Type: Multiple-choice

In which situation, parquet cudf is slower than csv cudf.



Papon CHOONHAKLAI's Response

skips_row=20,col_indexes = range(0, 7000)

Correct answer

skips_row=20,col_indexes = range(0, 7000)

Question 45 **Correct**

 20 secs  5 points

Question Type: Multiple-choice

In the visualization of DLI, what is the output of the callback? Select all that apply.

Papon CHOONHAKLAI's Response

Graph

Correct answer

Graph

Question 46 Correct

35 secs 5 points

Question Type: Multiple-choice

what is correct about the assessment? select all that apply.

Papon CHOONHAKLAI's Response

need to convert to pandas df before display on dash figure need to drop na value set na_value to dummy -9999

Correct answer

need to convert to pandas df before display on dash figure need to drop na value set na_value to dummy -9999

Question 47 Correct

28 secs 5 points

Question Type: Reorder

What is the typical process for workflow pipeline? order it.

Papon CHOONHAKLAI's Response

add new features
fillmissing
normalize

Correct answer

add new features
fillmissing
normalize

Question 48 Incorrect

32 secs 0 points

Question Type: Multiple-choice

What is needed for mqtt installation? select all that apply

Papon CHOONHAKLAI's Response

mosquitto-clients mosquitto_pub

Correct answer

mosquitto mosquitto-clients

Question 49 Correct

25 secs 1 points

Question Type: Multiple-choice

what is correct about MQTT? select all that apply.


Papon CHOONHAKLAI's Response


default port is 1883 '#' is used as wildcard multilevel,

Correct answer

default port is 1883 '#' is used as wildcard multilevel,

Question 50 **Correct**

 16 secs

 1 points

Question Type: Multiple-choice

MQTT can do..... and while Kafka model is producer/ consumer

Papon CHOONHAKLAI's Response

publish/subscribe

Correct answer

publish/subscribe